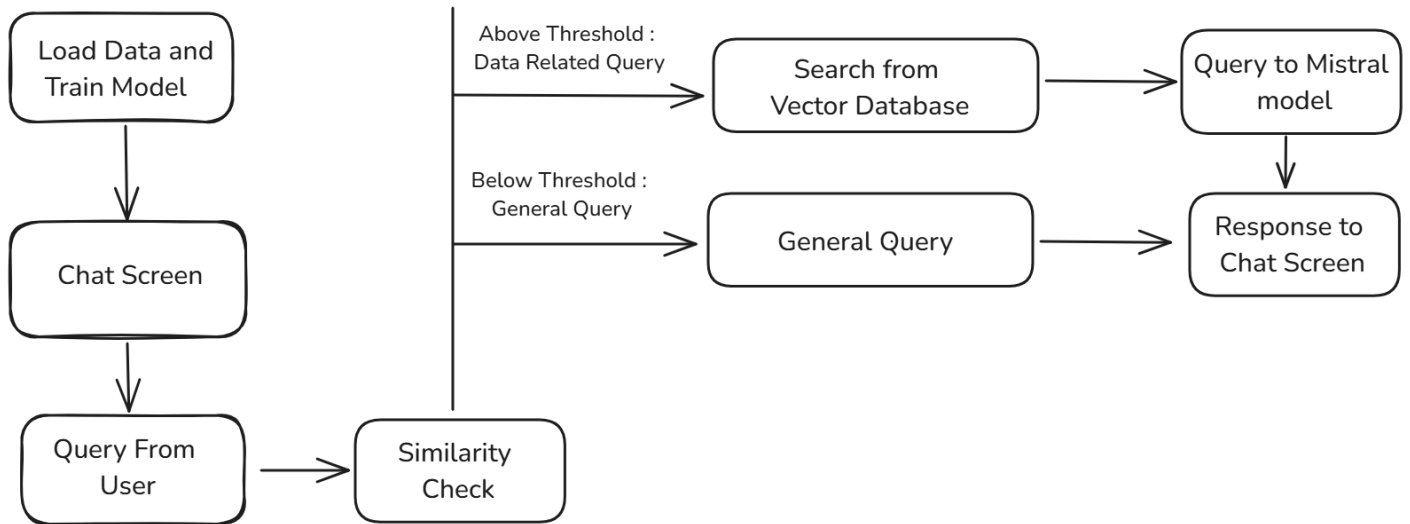


TECHNICAL DOCUMENTATION

System Architecture explanation



1. Load Data and Train Model

- The system begins by loading scraped restaurant/menu data and possibly embedding it using a model.
- A vector index (Pinecone) is prepared for retrieval tasks.

2. Chat Screen

- The frontend interface where users input their queries.

3. Query from User

- User submits a question—this could be about menus, pricing, ratings, or general queries.

4. Similarity Check

- The query is embedded and compared to vectors in the database using cosine similarity check.
- A similarity threshold is used to determine the nature of the query:
 - Above threshold → Considered data-related, requires retrieval.
 - Below threshold → Treated as a general query.

For Data-Related Queries (Above Threshold)

5. Search from Vector Database
 - Queries are matched against vector representations of restaurant/menu data to fetch relevant context.
6. Query to Mistral Model
 - The context + user query is sent to the Mistral language model (or equivalent) to generate a smart response.
7. Response to Chat Screen
 - The final answer is displayed to the user.

Implementation details and design decisions

Implementation Details

1. Framework: The chatbot is built using Flask, which serves both as the API layer and optionally the chat frontend.
2. Scraping Tool:
 - a. Selenium is used to scrape Zomato for restaurant and menu information.
 - b. The data is stored as CSV files, one for restaurant metadata and another for menu items.
3. Data Preprocessing:
 - a. The scraped data is cleaned, structured, and then converted into vector embeddings using *multi-qa-MiniLM-L6-cos-v1*.
 - b. Each menu item and restaurant is embedded individually to ensure fine-grained retrieval.
4. Vector Database:
 - a. Pinecone is used to store and search through high-dimensional vector representations of the menu and restaurant data.
5. Query Handling:
 - a. Each user query is embedded and compared against stored vectors.

- b. A similarity threshold is used to determine if the query is related to the restaurant data or is a general query.
- 6. Language Model:
 - a. Mistral-7B is used to generate context-aware responses when a relevant context is retrieved.
 - b. If the query doesn't require data context (falls below threshold), the model responds directly.
- 7. API Tokens:
 - a. Hugging Face API and Pinecone API keys are loaded into rag_engine.py.
- 8. Health Check:
 - a. A simple health check endpoint (/health-check) is available for verifying setup and deployment.

Design Decisions

1. Threshold-based Query Routing
 - This avoids unnecessary Pinecone hits for queries that don't relate to restaurant data, optimizing cost and latency.
2. Vector Embedding Granularity
 - Menu items and restaurant info are embedded separately so the bot can return detailed responses about dishes, pricing, and dietary tags.
3. Use of Open Models
 - Using Hugging Face and Mistral ensures the system is cost-efficient and open-source compatible along with specific amount of accuracy.
4. Lightweight UI with Flask
 - Keeps the frontend minimal while leaving room to scale into a richer UI with just HTML serving as frontend along with flask as backend.

Challenges faced and solutions implemented

1. Data Collection

Problem:

A lot of restaurants don't have their own websites.

What We Did:

We had to rely on the Zomato site to extract data like restaurant names, locations, contact details, menu items, and other useful info.

2. Model Selection

Problem:

There were many free models available on Hugging Face, so choosing the right one was tricky.

What We Did:

We tested 4–5 different models to compare their outputs and picked the one that worked best for our use case.

3. Setting the Right Similarity Threshold

Problem:

If the similarity score was too low, data that belonged together didn't get grouped. If it was too high, unrelated data got matched.

What We Did:

We fine-tuned the similarity threshold to make sure the right data was matched properly — neither too strict nor too loose.

Future Improvement opportunities

1. Focus more on the dining experience by adding extra context like:

- Real-time location data (how far the restaurant is)
- Traffic conditions (so users can estimate travel time)
- Weather info (because people might not want to go out when it's raining)

2. Handle Multi-Cuisine Restaurants Better

Some places serve multiple cuisines, which can confuse the tagging.

- Use NLP techniques like Named Entity Recognition (NER) to classify food types better
- Build a small cuisine ontology or food category tree for consistency

3. Include Price & Offers in Analysis

People often choose based on price or discounts.

- Extract and highlight average pricing per person
- Use OCR or regex parsing to detect offers from images or menus

4. Voice Search Support

A lot of people use voice for convenience.

- Add speech-to-text functionality for queries
- Use intent classification to understand what the user is

5. Make it Mobile-Friendly

Since most users are on phones:

- Use responsive UI design
- Add swipe-based filters for easier interaction